Partial Regularization of First-Order Resolution Proofs

Jan Gorzny[1] [*] and Bruno Woltzenlogel Paleo[2,3]

[1] `jgorzny@uvic.ca`, University of Victoria, Canada
[2] `bruno@logic.at`, Vienna University of Technology, Austria
[3] Australian National University

Abstract. This paper describes the generalization of the proof compression algorithm `RecyclePivotsWithIntersection` from propositional to first-order logic. The generalized algorithm performs partial regularization of resolution proofs containing resolution and factoring inferences with *unification*, as generated by many automated theorem provers. An empirical evaluation of the generalized algorithm and its combinations with `GreedyLinearFirstOrderLowerUnits` is also presented.

1 Introduction

First-order automated theorem provers, commonly based on resolution and superposition calculi, have recently achieved a high degree of maturity. Proof production is a key feature that has been gaining importance, since proofs are crucial for applications that require certification of a prover's answers or information extractable from proofs (e.g. unsat cores, interpolants, instances of quantified variables). Nevertheless, proof production is non-trivial [**?**], and the best, most efficient provers do not necessarily generate the best, least redundant proofs.

For proofs using propositional resolution generated by SAT- and SMT-solvers, there is a wide variety of proof compression techniques. Algebraic properties of the resolution operation that might be useful for compression were investigated in [5]. Compression algorithms based on rearranging and sharing chains of resolution inferences have been developed in [2] and [11]. Cotton [4] proposed an algorithm that compresses a refutation by repeatedly splitting it into a proof of a heuristically chosen literal ℓ and a proof of $\bar{\ell}$, and then resolving them to form a new refutation. The `Reduce&Reconstruct` algorithm [10] searches for locally redundant subproofs that can be rewritten into subproofs of stronger clauses and with fewer resolution steps. A linear time proof compression algorithm based on partial regularization was proposed in [3] and improved in [6].

In contrast, there has been much less work on simplifying first-order proofs. For tree-like sequent calculus proofs, algorithms based on cut-introduction [9, 8] have been proposed. However, converting a DAG-like resolution or superposition proof, as usually generated by current provers, into a tree-like sequent calculus proof may increase the size of the proof. For arbitrary proofs in the TPTP

---

[12] format (including DAG-like first-order resolution proofs), there is a simple algorithm [14] that looks for terms that occur often in any TSTP [12] proof and introduces abbreviations for these terms.

The work reported in this paper is part of a new trend that aims at lifting successful propositional proof compression algorithms to first-order logic. Our first target was the propositional `LowerUnits` algorithm, which delays resolution steps with unit clauses, and its lifting resulted in the `GreedyLinearFirstOrderLowerUnits` (`GFOLU`) algorithm [7]. Here we continue this line of research by lifting the `RecyclePivotsWithIntersection` (`RPI`) algorithm [6], which is an improvement of the `RecyclePivots` (`RP`) algorithm [3], providing better compression on proofs where nodes have several children.

Section 2 introduces the first-order resolution calculus and the notations used in this paper. Section 4 discusses the challenges that arise in the first-order case (mainly due to unification), which are not present the propositional case. Section 5 describes an algorithm that overcomes these challenges. Section 6 concludes the paper by presenting experimental results obtained by applying this algorithm, and also its combinations with `GFOLU`, on hundreds of proofs generated with the `SPASS` theorem prover.

2 The Resolution Calculus

We assume that there are infinitely many variable symbols (e.g. X, Y, Z, X_1, X_2, ...), constant symbols (e.g. a, b, c, a_1, a_2, ...), function symbols of every arity (e.g f, g, f_1, f_2, ...) and predicate symbols of every arity (e.g. p, q, p_1, p_2,...). A *term* is any variable, constant or the application of an n-ary function symbol to n terms. An *atomic formula* (*atom*) is the application of an n-ary predicate symbol to n terms. A *literal* is an atom or the negation of an atom. The *complement* of a literal ℓ is denoted $\bar{\ell}$ (i.e. for any atom p, $\bar{p} = \neg p$ and $\overline{\neg p} = p$). The set of all literals is denoted $\mathcal{L}$. A *clause* is a multiset of literals. $\bot$ denotes the *empty clause*. A *unit clause* is a clause with a single literal. Sequent notation is used for clauses (i.e. $p_1, \ldots, p_n \vdash q_1, \ldots, q_m$ denotes the clause $\{\neg p_1, \ldots, \neg p_n, q_1, \ldots, q_m\}$). $\mathrm{FV}(t)$ (resp. $\mathrm{FV}(\ell)$, $\mathrm{FV}(\Gamma)$) denotes the set of variables in the term t (resp. in the literal ℓ and in the clause Γ). A *substitution* $\{X_1 \backslash t_1, X_2 \backslash t_2, \ldots\}$ is a mapping from variables $\{X_1, X_2, \ldots\}$ to, respectively, terms $\{t_1, t_2, \ldots\}$. The application of a substitution σ to a term t, a literal ℓ or a clause Γ results in, respectively, the term $t\sigma$, the literal $\ell\sigma$ or the clause $\Gamma\sigma$, obtained from t, ℓ and Γ by replacing all occurrences of the variables in σ by the corresponding terms in σ. The set of all substitutions is denoted $\mathcal{S}$. A *unifier* of a set of literals is a substitution that makes all literals in the set equal. A *resolution proof* is a directed acyclic graph of clauses where the edges correspond to the inference rules of resolution and contraction (as explained in detail in Definition 1). A *resolution refutation* is a resolution proof with root $\bot$.

Definition 1 (First-Order Resolution Proof).
A directed acyclic graph $\langle V, E, \Gamma \rangle$, where V is a set of nodes and E is a set

of edges labeled by literals and substitutions (i.e. $E \subset V \times 2^{\mathcal{L}} \times \mathcal{S} \times V$ and $v_1 \xrightarrow[\sigma]{\ell} v_2$ denotes an edge from node v_1 to node v_2 labeled by the literal ℓ and the substitution σ), is a proof of a clause Γ iff it is inductively constructible according to the following cases:

- ***Axiom:*** *If Γ is a clause, $\widehat{\Gamma}$ denotes some proof $\langle \{v\}, \varnothing, \Gamma \rangle$, where v is a new (axiom) node.*
- ***Resolution:*** *If ψ_L is a proof $\langle V_L, E_L, \Gamma_L \rangle$ with $\ell_L \in \Gamma_L$ and ψ_R is a proof $\langle V_R, E_R, \Gamma_R \rangle$ with $\ell_R \in \Gamma_R$, and σ_L and σ_R are substitutions such that $\ell_L \sigma_L = \overline{\ell_R} \sigma_R$ and $\mathrm{FV}((\Gamma_L \setminus \{\ell_L\}) \sigma_L) \cap \mathrm{FV}((\Gamma_R \setminus \{\ell_R\}) \sigma_R) = \emptyset$, then $\psi_L \odot_{\ell_L \ell_R}^{\sigma_L \sigma_R} \psi_R$ denotes a proof $\langle V, E, \Gamma \rangle$ s.t.*

$$V = V_L \cup V_R \cup \{v\}$$

$$E = E_L \cup E_R \cup \left\{ \rho(\psi_L) \xrightarrow[\sigma_L]{\{\ell_L\}} v, \rho(\psi_R) \xrightarrow[\sigma_R]{\{\ell_R\}} v \right\}$$

$$\Gamma = (\Gamma_L \setminus \{\ell_L\}) \sigma_L \cup (\Gamma_R \setminus \{\ell_R\}) \sigma_R$$

where v is a new (resolution) node and $\rho(\varphi)$ denotes the root node of φ. The resolved atom ℓ is such that $\ell = \ell_L \sigma_L = \overline{\ell_R} \sigma_R$ or $\ell = \overline{\ell_L} \sigma_L = \ell_R \sigma_R$.
- ***Contraction:*** *If ψ' is a proof $\langle V', E', \Gamma' \rangle$ and σ is a unifier of $\{\ell_1, \ldots \ell_n\}$ with $\{\ell_1, \ldots \ell_n\} \subseteq \Gamma'$, then $\lfloor \psi \rfloor_{\{\ell_1, \ldots \ell_n\}}^{\sigma}$ denotes a proof $\langle V, E, \Gamma \rangle$ s.t.*

$$V = V' \cup \{v\}$$

$$E = E' \cup \{ \rho(\psi') \xrightarrow[\sigma]{\{\ell_1, \ldots \ell_n\}} v \}$$

$$\Gamma = (\Gamma' \setminus \{\ell_1, \ldots \ell_n\}) \sigma \cup \{\ell\}$$

where v is a new (contraction) node, $\ell = \ell_k \sigma$ (for any $k \in \{1, \ldots, n\}$) and $\rho(\varphi)$ denotes the root node of φ. □

3   The Propositional Algorithm

RPI (formally defined in Appendix A) removes *irregularities*, which are resolution inferences with a node η when the resolved literal (a.k.a. *pivot*) occurs as the pivot of another inference located below in the path from η to the root of the proof. In the worst case, regular resolution proofs can be exponentially bigger than irregular ones, but RPI takes care of regularizing the proof only partially, removing inferences only when this does not enlarge the proof.

RPI traverses the a proof twice from the bottom-up. On the first traversal, it stores for each node a set of *safe literals* that get resolved in all paths below it in the proof, or which already occurred in the root clause of the proof). If one of the node's resolved literals belongs to the set of safe literals, then it is possible to *regularize* the node by replacing it by one of its parents. Such a node is marked as deleted. On the second traversal, regularization is performed: nodes that have a parent node marked deleted are replaced.

The `RPI` and the `RP` algorithms differ from each other mainly in the computation of the safe literals of a node that has many children. While the former returns the intersection as shown in Algorithm 6, the latter returns the empty set. Further, while in `RPI` the safe literals of the root node contain all the literals of the root clause, in `RP` the root node is always assigned an empty set of literals.

4 First-Order Challenges

In this section, we describe challenges that have to be overcome in order to successfully adapt `RPI` to the first-order case. The first example illustrates the need to take unification into account. The other two examples discuss complex issues that can arise when unification is taken into account in a naive way.

Example 1. Consider the following irregular proof ψ. Naively computed, the safe literals for η_3 are $\{\vdash q(c),\ p(a, X)\}$. η_3's left pivot $p(W, X) \in \eta_1$ is unifiable with $p(a, X)$ in its safe literals, and thus the proof can be regularized by recycling η_1.

$$
\cfrac{\cfrac{\eta_1\colon\ \vdash p(W, X) \qquad \eta_2\colon p(W, X)\ \vdash q(c)}{\eta_3\colon\ \vdash q(c)} \qquad \cfrac{}{\eta_4\colon q(c)\ \vdash p(a, X)}}{\cfrac{\eta_5\colon\ \vdash p(a, X) \qquad\qquad \eta_6\colon p(Y, b)\ \vdash}{\psi\colon \bot}}
$$

Regularization of the proof by recycling η_1 results in removing the inference between η_2 and η_3, which in turn replaces η_3 by η_1. Since η_1 cannot be resolved against η_4, and η_1 contained safe literals, η_5 is replaced by η_1. The result is the much shorter proof below.

$$
\cfrac{\eta_1\colon\ \vdash p(W, X) \qquad \eta_6\colon p(Y, b) \vdash}{\psi'\colon \bot}
$$

Unlike in the propositional case, where the pivots and their corresponding safe literal list are all syntactically equal, in the first-order case, this is not necessarily the case. As illustrated above, $p(W, X)$ and $p(a, X)$ are not syntactically equal. Nevertheless, they are unifiable, and the proof can be regularized.

Example 2. There are cases, as shown below, that require more careful care when attempting to regularize. Again, naively computed, the safe literals for η_3 are $\{\vdash q(c),\ p(a, X)\}$, and so η_1 appears to be a candidate for regularization.

$$
\cfrac{\cfrac{\eta_1\colon\ \vdash p(a, c) \qquad \eta_2\colon p(a, c)\ \vdash q(c)}{\eta_3\colon\ \vdash q(c)} \qquad \cfrac{}{\eta_4\colon q(c)\ \vdash p(a, X)}}{\cfrac{\eta_5\colon\ \vdash p(a, X) \qquad\qquad \eta_6\colon p(Y, b)\ \vdash}{\psi\colon \bot}}
$$

However, if we attempt to regularize the proof, the same series of actions as in Example 1 would require resolution between η_1 and η_5, which is not possible.

Thus we require the following property be satisfied before regularization.

Definition 2. *Let η be a clause with literal ℓ' with corresponding safe literal ℓ which is resolved against literals ℓ_1, ..., ℓ_n in a proof ψ. η is said to satisfy the pre-regularization unifiability property in ψ if $\ell_1,\ldots,\ell_n$, and $\overline{\ell'}$ are unifiable.*

One technique to ensure this property is met is to apply the unifier of a resolution to each resolvent before computing the safe literals. In the case of Example 2, this would result in η_3 having the safe literals $\{\vdash q(c),\ p(a,b)\}$, and now it is clear that the literal in η_1 is not safe.

Example 3. Satisfying the pre-regularization unifiability property is not sufficient to attempt regularization. Consider the proof ψ below. After collecting the safe literals, η_3's safe literals are $\{q(T,V), p(c,d) \vdash q(f(a,e),c)\}$.

$$
\dfrac{
\dfrac{\eta_1\colon p(U,V) \vdash q(f(a,V),U) \qquad \dfrac{\eta_2\colon q(f(a,X),Y), q(T,X) \vdash q(f(a,Z),Y)}{\dfrac{\eta_3\colon p(U,V), Q(T,V) \vdash q(f(a,Z),U)\ \eta_4\colon\ \vdash q(R,S)}{\eta_5\colon p(U,V) \vdash q(f(a,Z),U)}}}{\dfrac{\eta_6\colon\ \vdash p(c,d) \qquad \eta_7\colon\ \vdash q(f(a,Z),c)}{\eta_8\colon q(f(a,e),c) \vdash}}}{\psi\colon \bot}
$$

Since $q(f(a,X),Y) \in \eta_2$ and $q(T,V)$ (in η_3's safe literals) are unifiable, regularization would be attempted. In this case, the inference between η_2 and η_3 for would be removed, and as a result, η_3 will be replaced with η_1. η_1 does not contain the required pivot for η_5, and so η_5 is also replaced with η_1, and resolution is attempted before η_1 and η_6, which results in η_7', and an inability to complete the proof, as shown below.

$$
\dfrac{\eta_8\colon Q(f(a,e)c) \vdash \qquad \dfrac{\eta_6\colon\ \vdash P(c,d) \qquad \eta_1\colon P(U,V) \vdash Q(f(a,V)U)}{\eta_7'\colon\ \vdash Q(f(a,d)c)}}{\psi'\colon\ ??}
$$

In order to avoid these scenarios, we perform an additional check during inference removal. The node $\eta*$ which will replace a resolution η (because η would have a deleted parent), must be entirely contained, via unification which modifies only $\eta*$'s variables, in the safe literals of η. In this example, η_1 does not satisfy this property. This check is not necessary in the propositional case, as the replacement node would be contained exactly in the set of safe literals, and would not change lower in the proof.

5 First-Order RecyclePivotsWithIntersection

This section presents `FirstOrderRecyclePivotsWithIntersection` (FORPI), Algorithm 1, a first order generalization of `RPI`. FORPI traverses the proof in a bottom-up manner, storing for every node a set of safe literals. If one of the node's resolved literals can be unified to a literal in the set of safe literals, then it may be possible to regularize the node by replacing it by one of its parents.

In the propositional case, regularization of a node replaces it by the parent whose clause contains the resolved literal that is safe. In the first order case, because unification introduces complications like those seen in Example 3, we

```
    input  : A first-order proof ψ
    output: A possibly less-irregular first-order proof ψ′
1  ψ′ ← ψ;
2  traverse ψ′ bottom-up and foreach node η in ψ′ do
3      if η is a resolvent node then
4          setSafeLiterals(η) ;
5          regularizeIfPossible(η)
6  ψ′ ← fix(ψ′) ;
7  return ψ′;
```

Algorithm 1: `FORPI`

```
    input  : A node ψ = ψ_L ⊙_{ℓ_L ℓ_R}^{σ_L σ_R} ψ_R
    output: nothing (but the proof containing ψ may be changed)
1  if ∃σ and ℓ ∈ ψ.safeLiterals such that σℓ = ℓ_R or ℓ = σℓ_R then
2      if ∃σ′ such that σ′ψ_R ⊆ ψ.safeLiterals then
3          replace ψ_L by deletedNodeMarker ;
4          mark ψ as regularized
5  else if ∃σ and ℓ ∈ ψ.safeLiterals such that σℓ = ℓ_L or ℓ = σℓ_L then
6      if ∃σ′ such that σ′ψ_L ⊆ ψ.safeLiterals then
7          replace ψ_R by deletedNodeMarker ;
8          mark ψ as regularized
```

Algorithm 2: `regularizeIfPossible`

ensure that the replacement parent is (possibly after unification) contained entirely in the safe literals. This ensures that the remainder of the proof does not expect a variable to be unified to different values simultaneously. Similar to `RPI`, instead of replacing the irregular node by one of its parents immediately, its other parent is replaced by `deletedNodeMarker`, as shown in Algorithm 5. As in the propositional case, fixing of the proof is postponed to another (single) traversal, as regularization proceeds bottom up and only nodes below a regularized node may require fixing. During fixing, the irregular node is actually replaced by the parent that is not `deletedNodeMarker`. With careful bookkeeping, it is often possible to contract nodes during proof fixing to compress the proof further.

Changing between generalizations of `RPI` and `RP` is easily accomplished in the first order case by changing lines 11 and 2, respectively, of Algorithm 6. This makes a difference only when the proof is not a refutation.

The set of safe literals for a node ψ is computed from the set of safe literals of its children (cf. Algorithm 6), similar to the propositional case, but additionally applies unifiers to the resolved pivots (cf. Example 2).

> **input** : A first order resolution node ψ
> **output**: nothing (but the node ψ gets a set of safe literals)

1 **if** *ψ is a root node with no children* **then**
2 ψ.safeLiterals $\leftarrow \psi$.clause
3 **else**
4   **foreach** $\psi' \in \psi$.children **do**
5   **if** *ψ' is marked as regularized* **then**
6   safeLiteralsFrom$(\psi') \leftarrow \psi'$.safeLiterals ;
7   **else if** *$\psi' = \psi \odot_{\ell_L \ell_R}^{\sigma_L \sigma_R} \psi_R$ for some ψ_R* **then**
8   safeLiteralsFrom$(\psi') \leftarrow \psi'$.safeLiterals $\cup \{ \sigma_R \ell_R \}$
9   **else if** *$\psi' = \psi_L \odot_{\ell_L \ell_R}^{\sigma_L \sigma_R} \psi$ for some ψ_L* **then**
10   safeLiteralsFrom$(\psi') \leftarrow \psi'$.safeLiterals $\cup \{ \sigma_L \ell_L \}$
11 ψ.safeLiterals $\leftarrow \bigcap_{\psi' \in \psi.\text{children}}$ safeLiteralsFrom(ψ')

Algorithm 3: `setSafeLiterals`

6 Experiments

A prototype[1] version of `FORPI` has been implemented in the functional programming language Scala as part of the Skeptik library. Evaluation of this algorithm was performed on the same 308 real first-order proofs generated to evaluate `GFOLU`, and for consistency, the same system and metrics were used (see [7]).

Figure 1 (a) shows the compression results of applying `FORPI` and `GFOLU` to the same proof (in both application orders), as well as each of these algorithms individually. Unsurprisingly, applying both algorithms generally does better than either algorithm alone. On this data set, `FORPI` compresses provides some compression to the original proofs, `GFOLU` is responsible for most of the compression. `FORPI` compresses only three proofs already compressed by `GFOLU`. `RPI` performs best when the proofs are tall; `FORPI` will likely perform similarly. However, the proofs in this data set are relatively short, and those compressed by `GFOLU` first are even shorter. Thus, the performance of `FORPI` is not unsurprising.

Figure 1 (b) shows that the order of compression may matter less than in than in the propositional case, although more data is needed to confirm. The number of points above and below the main diagonal are the same; however, the points below may simply be the result of `GFOLU` being more likely to compress such short proofs. If so, that would imply that Running `FORPI` after `GFOLU` is more successful, which would be consistent with propositional results for these algorithms.

SPASS required approximately 40 minutes to solve and generate the proofs; the total time for `GFOLU` and `FORPI` to be executed on all 308 proofs was just under 8 seconds (both include parsing time). These compression algorithms continue to be very fast, and may simplify the proof considerably for a relatively quick time cost.

---

[1] Source code available at https://github.com/jgorzny/Skeptik

(a) Compressed length against input length (b) FORPI (GFOLU (p)) vs. GFOLU (FORPI (p))

Fig. 1: Experimental results

References

1. *Logic for Programming, Artificial Intelligence, and Reasoning - 16th Intl. Conf., Dakar, Senegal, Rev. Selected Papers*, LNCS. Springer, 2010.
2. Hasan Amjad. Compressing propositional refutations. *Electr. Notes Theor. Comput. Sci.*, 185:3–15, 2007.
3. O. Bar-Ilan, O. Fuhrmann, S. Hoory, O. Shacham, and O. Strichman. Linear-time reductions of resolution proofs. In *Haifa Verif. Conf.*, LNCS, pages 114–128. Springer, 2008.
4. Scott Cotton. Two techniques for minimizing resolution proofs. In Ofer Strichman and Stefan Szeider, editors, *Theory and Applications of Satisfiability Testing SAT 2010*, volume 6175 of *LNCS*, pages 306–312. Springer, 2010.
5. P. Fontaine, S. Merz, and B. Woltzenlogel Paleo. Exploring and exploiting algebraic and graphical properties of resolution. In *8th Intl. Wkshp. on SMT*, Edinburgh, 2010.
6. P. Fontaine, S. Merz, and B. Woltzenlogel Paleo. Compression of propositional resolution proofs via partial regularization. In *CADE*, LNCS, pages 237–251. Springer, 2011.
7. J. Gorzny and B. Woltzenlogel Paleo. Towards the compression of first-order resolution proofs by lowering unit clauses. In *CADE*, 2015.
8. S. Hetzl, A. Leitsch, G. Reis, and D. Weller. Algorithmic introduction of quantified cuts. *Theor. Comput. Sci.*, 549:1–16, 2014.
9. B. Woltzenlogel Paleo. Atomic cut introduction by resolution: Proof structuring and compression. In *LPAR-16* [1], pages 463–480.
10. S. F. Rollini, R. Bruttomesso, and N. Sharygina. An efficient and flexible approach to resolution proof reduction. In *Hardware and Software: Verification and Testing*, LNCS, pages 182–196. Springer, 2011.

```
   input  : A proof ψ
   output: A possibly less-irregular proof ψ′
 1 ψ′ ← ψ;
 2 traverse ψ′ bottom-up and foreach node η in ψ′ do
 3     if η is a resolvent node then
 4         setSafeLiterals(η) ;
 5         regularizeIfPossible(η)
 6 ψ′ ← fix(ψ′) ;
 7 return ψ′;
```

Algorithm 4: `RPI`

11. Carsten Sinz. Compressing propositional proofs by common subproof extraction.
    In Roberto Moreno-Díaz, Franz Pichler, and Alexis Quesada-Arencibia, editors,
    EUROCAST, volume 4739 of *LNCS*, pages 547–555. Springer, 2007.
12. G. Sutcliffe. The TPTP Problem Library and Associated Infrastructure: The FOF
    and CNF Parts, v3.5.0. *Journal of Automated Reasoning*, 43(4):337–362, 2009.
13. G. S. Tseitin. On the complexity of derivation in propositional calculus. In J. Siek-
    mann and G. Wrightson, editors, *Automation of Reasoning: Classical Papers in
    Computational Logic 1967-1970*, volume 2. Springer-Verlag, 1983.
14. J. Vyskocil, D. Stanovský, and J. Urban. Automated proof compression by inven-
    tion of new definitions. In *LPAR-16* [1], pages 447–462.

A   Algorithm `RecyclePivotsWithIntersection`

`RecyclePivotsWithIntersection` (`RPI`) aims at compressing irregular proofs.
It can be seen as a simple but significant modification of the `RP` algorithm de-
scribed in [3], from which it derives its name. Although in the worst case full
regularization can increase the proof length exponentially [13], these algorithms
show that many irregular proofs can have their length decreased if a careful
partial regularization is performed.

Consider an irregular proof of the form $\psi[\eta \odot_p \psi'[\eta' \odot_p \eta'']]$ and assume,
without loss of generality, that $p \in \eta$ and $p \in \eta'$. Then, if $\eta' \odot_p \eta''$ is replaced
by η'' within the proof-context $\psi'[\ ]$, the clause $\eta \odot_p \psi'[\eta'']$ subsumes the clause
$\eta \odot_p \psi'[\eta' \odot_p \eta'']$, because even though the literal $\neg p$ of η'' is propagated down, it
gets resolved against the literal p of η later on below in the proof. More precisely,
even though it might be the case that $\neg p \in \psi'[\eta'']$ while $\neg p \notin \psi'[\eta' \odot_p \eta'']$, it is
necessarily the case that $\neg p \notin \eta \odot_p \psi'[\eta' \odot_p \eta'']$ and $\neg p \notin \eta \odot_p \psi'[\eta'']$.

Although the remarks above suggest that it is safe to replace $\eta' \odot_p \eta''$ by
η'' within the proof-context $\psi'[\ ]$, this is not always the case. If a node in $\psi'[\ ]$
has a child in $\psi[\ ]$, then the literal $\neg p$ might be propagated down to the root
of the proof, and hence, the clause $\psi[\eta \odot_p \psi'[\eta'']]$ might not subsume the clause
$\psi[\eta \odot_p \psi'[\eta' \odot_p \eta'']]$. Therefore, it is only safe to do the replacement if the literal
$\neg p$ gets resolved in all paths from η'' to the root or if it already occurs in the
root clause of the original proof $\psi[\eta \odot_p \psi'[\eta' \odot_p \eta'']]$.

These observations lead to the idea of traversing the proof in a bottom-up manner, storing for every node a set of *safe literals* that get resolved in all paths below it in the proof (or that already occurred in the root clause of the original proof). Moreover, if one of the node's resolved literals belongs to the set of safe literals, then it is possible to regularize the node by replacing it by one of its parents (cf. Algorithm 4).

The regularization of a node should replace a node by one of its parents, and more precisely by the parent whose clause contains the resolved literal that is safe. After regularization, all nodes below the regularized node may have to be fixed. However, since the regularization is done with a bottom-up traversal, and only nodes below the regularized node need to be fixed, it is again possible to postpone fixing and do it with only a single traversal afterwards. Therefore, instead of replacing the irregular node by one of its parents immediately, its other parent is replaced by `deletedNodeMarker`, as shown in Algorithm 5. Only later during fixing, the irregular node is actually replaced by its surviving parent (i.e. the parent that is not `deletedNodeMarker`).

The set of safe literals of a node η can be computed from the set of safe literals of its children (cf. Algorithm 6). In the case when η has a single child ς, the safe literals of η are simply the safe literals of ς together with the resolved literal p of ς belonging to η (p is safe for η, because whenever p is propagated down the proof through η, p gets resolved in ς). It is important to note, however, that if ς has been marked as regularized, it will eventually be replaced by η, and hence p should not be added to the safe literals of η. In this case, the safe literals of η should be exactly the same as the safe literals of ς. When η has several children, the safe literals of η w.r.t. a child ς_i contain literals that are safe on all paths that go from η through ς_i to the root. For a literal to be safe for all paths from η to the root, it should therefore be in the intersection of the sets of safe literals w.r.t. each child.

The `RP` and the `RPI` algorithms differ from each other mainly in the computation of the safe literals of a node that has many children. While `RPI` returns the intersection as shown in Algorithm 6, `RP` returns the empty set (cf. Algorithm 7). Additionally, while in `RPI` the safe literals of the root node contain all the literals of the root clause, in `RP` the root node is always assigned an empty set of literals. (Of course, this makes a difference only when the proof is not a refutation.) Note that during a traversal of the proof, the lines from 5 to 10 in Algorithm 6 are executed as many times as the number of edges in the proof. Since every node has at most two parents, the number of edges is at most twice the number of nodes. Therefore, during a traversal of a proof with n nodes, lines from 5 to 10 are executed at most $2n$ times, and the algorithm remains linear. In our prototype implementation, the sets of safe literals are instances of Scala's `mutable.HashSet` class. Being mutable, new elements can be added efficiently. And being HashSets, membership checking is done in constant time in the average case, and set intersection (line 12) can be done in $O(k.s)$, where k is the number of sets and s is the size of the smallest set.

<div style="border:1px solid">

input : A node η

output: nothing (but the proof containing η may be changed)

1 **if** η.rightResolvedLiteral $\in$ η.safeLiterals **then**

2   replace left parent of η by `deletedNodeMarker` ;

3   mark η as regularized

4 **else if** η.leftResolvedLiteral $\in$ η.safeLiterals **then**

5   replace right parent of η by `deletedNodeMarker` ;

6   mark η as regularized

</div>

Algorithm 5: `regularizeIfPossible`

<div style="border:1px solid">

input : A node η

output: nothing (but the node η gets a set of safe literals)

1 **if** *η is a root node with no children* **then**

2 η.safeLiterals $\leftarrow$ η.clause

3 **else**

4   **foreach** $\eta' \in \eta$.children **do**

5   **if** *η' is marked as regularized* **then**

6   safeLiteralsFrom(η') $\leftarrow$ η'.safeLiterals ;

7   **else if** *η is left parent of η'* **then**

8   safeLiteralsFrom(η') $\leftarrow$ η'.safeLiterals $\cup$ { η'.rightResolvedLiteral } ;

9   **else if** *η is right parent of η'* **then**

10   safeLiteralsFrom(η') $\leftarrow$ η'.safeLiterals $\cup$ { η'.leftResolvedLiteral } ;

11 η.safeLiterals $\leftarrow$ $\bigcap_{\eta' \in \eta.\text{children}}$ safeLiteralsFrom(η')

</div>

Algorithm 6: `setSafeLiterals`

<div style="border:1px solid">

input : A node η

output: nothing (but the node η gets a set of safe literals)

1 **if** *η is a root node with no children* **then**

2 η.safeLiterals $\leftarrow$ $\emptyset$

3 **else**

4   **if** *η has only one child η'* **then**

5   **if** *η' is marked as regularized* **then**

6 η.safeLiterals $\leftarrow$ η'.safeLiterals ;

7   **else if** *η is left parent of η'* **then**

8 η.safeLiterals $\leftarrow$ η'.safeLiterals $\cup$ { η'.rightResolvedLiteral } ;

9   **else if** *η is right parent of η'* **then**

10 η.safeLiterals $\leftarrow$ η'.safeLiterals $\cup$ { η'.leftResolvedLiteral } ;

11   **else**

12 η.safeLiterals $\leftarrow$ $\emptyset$

</div>

Algorithm 7: `setSafeLiterals` for RP